

Design Exercise 1

Engineering Notebook

Adarsh Hiremath and Andrew Palacci

February 19, 2023

Introduction

In our design exercise, we implemented an end-to-end client-server chat application, first with our own wire protocol and later with gRPC. We built our chat application using Python with a simple tutorial and eventually built functionality on top of our minimal chat room application to fulfill the assignment specification. As currently implemented, our chat application supports the following:

- Account creation, login, deletion, and deactivation.
- Sending messages between accounts, even when some accounts are offline.
- Listing or filtering all users who have created accounts.

The source code for our chat application can be found [here](#). Throughout this notebook, Andrew and I intend to describe various design decisions made throughout the development of this project.

Programming Language

Initially, Andrew and I considered 3 programming languages to implement our client-server chat application in: Python, Java, and C++. While Professor Waldo recommended that we complete the project in Java, Andrew and I decided to use another language since we couldn't find an elegant solution for handling different versions of the Java Virtual Machine (JVM) during peer grading.

Ultimately, Andrew and I decided to implement our design exercise in Python because of the extensive documentation for the socket module. While C++ was the clear choice for low-level optimizations, Andrew and I wanted to pick a language that would be most beneficial for learning. After some exploration on the internet, it became clear that tutorials for socket programming in Python were far more accessible to us than tutorials for socket programming in C++.

The final benefit to using Python instead of C++ was code reproducibility. Most if not all students have Python natively on their machines. Any student can run our program with a simple command

allowing them to mirror the dependencies listed in a requirements.txt file, regardless of their machine. Conversely, maintaining code reproducibility across machines for a C++ implementation would require creating precompiled binaries for students, which we found to be far more tedious.

Server vs. Client

Before implementing our client or server, Andrew and I spent some time considering what input processing we wanted to put on the server side versus the client side. As will be explained later in this writeup, our wire protocol handles information differently depending on the operation code provided before a client types in the delimiter. Our options were to either process the operations on the client side, rejecting invalid inputs before even pinging the server, or handling the operations on the server side.

Andrew and I decided to process information sent by the client on the server side for a variety of reasons:

- A client side approach to input validation requires all clients to download updates if any operation specifications change. This results in versioning issues.
- Validation done on the client side results in code redundancy since every client has duplicate code specifying what inputs are allowed.
- Clients having access to the type of operations they can use to ping the server results in security issues.

Wire Protocol

Here, we will specify the functionalities supported by our wire protocol as well as provide further explanation about how our protocol was implemented. Our wire protocol utilizes a delimiter system to distinguish between different requests from the user. Our delimiter is '|', so any requests made to the server must include the appropriate command followed by the delimiter followed by any relevant parameters. For example, sending a message to another user would look something like:

's|<recipientUsername>|<message>.' Below, we've included a complete list of the arguments accepted by our wire protocol.

- Create an account. Usage: c|<username>
- Log into an account. Usage: l|<username>
- Send message.
Usage: s|<recipientUsername>|<message>
- List all users. Usage: u
- Filter accounts. Usage: f|<filterRegex>
- Delete account. Usage: d|<confirmUsername>

Commands like logging out, listing all users, and deleting an account become easier to implement with a delimiter system. Without a delimiter system, a client sending "d," for example, would be ambiguous for the server because there is no way to distinguish between sending the character "d" to another client or deleting an account. A delimiter system solves this because the client can still send the character "d" to a user as long as they specify their desired operation is a message send.

We use Python's `Lib/struct.py` to encode and decode messages efficiently using UTF-8. The two encoding and decoding schemes we considered were ASCII and UTF-8. At first, we considered implementing our encode / decode system using ASCII, since the size of any individual ASCII character is exactly one byte. This would make implementing character limits relatively easy to implement, since our argument to `connection.recv()` would be the maximum number of characters we want a user to be able to send using our chat application. However, since ASCII only supports A-Z, 0-9, and dashes, we decided to use a UTF-8 encoding system to allow users to use punctuation when necessary. We decided on a limit of 4096 bytes, which allows for approximately 1000 characters to be sent over our socket in the worst case scenario.

After receiving our bytes over the web socket and decoding, we do a simple `split('|')` operation to distinguish between the delimiter and arguments sent by the client. We decided to deal with white space using a simple `strip()` operation on every element in the split version of the client's command. This way, any spaces before or after the delimiter would be irrelevant and the client can still make commands as long as they follow the outline described above.

To keep our wire protocol simple and minimize code repetition, we decided to abstract away the implementations using helper functions for each operation. For example, this is how our code is structured for creating an account, logging in, and listing users. The code for the other operations follows a similar structure. Each operation code corresponds to a different helper function and each helper function

behaves differently depending on the specification of the operation.

```
if op_code == 'c':
    msg = create_account(msg_list[1],
                        connection)
elif op_code == 'l':
    msg = login(msg_list[1], connection)
elif op_code == 'u':
    msg = list_accounts()
```

Additionally, implementing our wire protocol using helper functions allowed us to easily make our server side send information to the user or print information in the console when applicable. In this case, our helper functions return information to print on the server and actually send information to the appropriate client within the function body. Here is an example of our implementation for our function that sends messages.

```
def send_msg(uuid, msg):
    print(f"connection: {live_users[uuid]}")
    live_users[uuid].send
    (msg.encode('UTF-8'))
    return f"\nMessage sent to {uuid}.\n"
```

Account Handling

To implement our account handling schema, we debated between using a database or simple Python dictionaries and lists. After consulting Professor Waldo, we decided to use a simple dictionary and list implementation since we assumed that there would be no Byzantine failures causing the server to unexpectedly shut down.

We have three main data structures that we use throughout our chat application.

- *pending_messages*: a dictionary with UUIDs as keys and pending messages as values.
- *accounts*: a dictionary with UUIDs as keys and account names as values that we call
- *live_users*: a dictionary with UUIDs as keys and connection references as values.

We decided to use a single dictionary for pending messages to handle the scenario where a user is offline. Storing all messages sent and received in a dictionary would be a waste of space because our server only cares if the user has seen the message intended for them. When a user goes online, our program checks whether there are any pending messages to be sent to them and sends messages accordingly. We have a single dictionary for accounts and their corresponding names in order to implement the listing users portion of the assignment. Finally, our live users dictionary stores all the users that are currently logged in so that the server can check if a user

is logged in before sending a message to them. We also store the references to the socket connection in order for the server to send messages over the right connection.

Original Wire Protocol vs. gRPC